

Direct device access from the SmartNIC towards datacenter disaggregation

Master's thesis meeting : week 7

Nicolas Jeanmenne

Table of contents

- State of the art
- SOTA reviews
- How do Scalio does NVMe-oF in the ARM cores ?
- Why does TCP-IO needs a smartNIC ?
- Conclusion
- TODOs for week 9

State of the art

- Started writing SOTA about why disaggregated architecture matters, as well as challenges that arise with these patterns and contributions
- Working on linking papers innovations by core concepts as asked
- Let's take a look at the \LaTeX

SOTA reviews

L^AT_EX sources (and PDF) will be available for reviews through the year on my [github repo](#)



How do Scalio does NVMe-oF in the ARM cores ?

- With standard NVMe-OF ARM cores are heavily solicited
 - Cores handle requests through a kernel module / userspace (via SPDK) ⇒ many cycles consumed
 - Becomes a bottleneck with many SSDs
- Scalio offload NVMe commands to the host channel adapter hardware
 - Clients read an in-memory hash table containing SSD index (the RDMA call discussed in previous meeting)
 - If cache miss then NVMe-OF TO call from the SSD




Figure 7: The workflows comparision between LEED's and Scalio.

Why does TCP-IO needs a smartNIC ?

- SmartNIC are well suited for the job, but authors state that any programmable device that can perform direct I/O for disks and network can be used.
 - Convenience use for P2PDMA and network I/O
- Highly optimized CPU-only approach can even be as performant (like we saw with ATLAS server in week 3)
- SmartNIC still outperforms classic Linux TCP and ATLAS for TLS

Conclusion

- SmartNIC may be not as much important than other types of accelerators
- Big picture for disaggregation becomes more clear

TODOs for week 9

- Investigate SPDK to write data blocks from the smartNIC towards NVMe on *Frodo*
- Reproduce experience from *Scalio*
- Find how important is the smartNIC related to the thesis
- Investigate NVMe Computational SSD (it is interesting for our storage disaggregation ?)

That's all for today !